

✓ Linear regression [9 pts]

In this homework, you will implement solution algorithms for linear regression.

Import libraries

Let's begin by importing some libraries.

```
print(__doc__)
import matplotlib.pyplot as plt
import numpy as np
from sklearn import datasets
%matplotlib inline
```

Automatically created module for IPython interactive environment

✓ Load dataset

Now, we are importing a dataset of diabetes. You can check the details on this dataset here:

https://scikit-learn.org/stable/datasets/toy_dataset.html#diabetes-dataset.

The dataset consists of 442 observations with 10 attributes (X) that may affect the progression of diabetes (y). Ten baseline variables, age, sex, body mass index, average blood pressure, and six blood serum measurements were obtained for each of $n = 442$ diabetes patients, as well as the response of interest, a quantitative measure of disease progression one year after baseline.

```
# Load the diabetes dataset
diabetes_X, diabetes_y = datasets.load_diabetes(return_X_y=True)
print('The shape of the input features:',diabetes_X.shape)
print('The shape of the output variable:',diabetes_y.shape)
```

The shape of the input features: (442, 10)
The shape of the output variable: (442,)

We will choose just one attribute from the ten attributes as an input variable.

```
# Use only one feature  
diabetes_X_one = diabetes_X[:, np.newaxis,2]  
print(diabetes_X_one.shape)  
  
(442, 1)
```

❖ Dataset split

Now, we split the dataset into two parts: training set and test set.

- training set: 422 samples
- test set: 20 samples

```
# Split the data into training/testing sets  
diabetes_X_train = diabetes_X_one[:-20]  
diabetes_X_test = diabetes_X_one[-20:]  
  
# Split the targets into training/testing sets  
diabetes_y_train = diabetes_y[:-20]  
diabetes_y_test = diabetes_y[-20:]  
  
print('Training input variable shape:', diabetes_X_train.shape)  
print('Test input variable shape:', diabetes_X_test.shape)
```

Training input variable shape: (422, 1)
Test input variable shape: (20, 1)

Linear regression

Assume that we have a hypothesis

$$h_{\theta}(x) = \theta_0 + \theta_1 x.$$

Your tasks:

- [3pts] implement {your own version} of the method of least-squares, compute and report θ_0 and θ_1 that minimize the residual sum of squares,)

$$\sum_{i=1}^N \frac{1}{2} (y^{(i)} - h_{\theta}(x^{(i)}))^2$$

[IMPORTANT] Do not just call the least square function from libraries, for example, `scipy.optimize.least_squares` from `scipy`. Doing so will result in 0 point. Using helping functions such as `numpy.linalg.inv` is okay.

- [3pts] implement your own version of the gradient descent algorithm, compute and report θ_0 and θ_1 that minimize the mean squared error

$$\sum_{i=1}^N \frac{1}{N} (y^{(i)} - h_{\theta}(x^{(i)}))^2$$

[NOTE] Notice that the loss function is mean-squared error.

- [3pts] derive the analytical expression of the gradient if the loss is defined as

$$\sum_{i=1}^N \frac{1}{2} (y^{(i)} - h_{\theta}(x^{(i)}))^2 + \frac{\lambda}{2} \|\theta\|_2^2,$$

where $\theta = [\theta_0, \theta_1]^T$

To check whether your computation is correct, consider using an API such as Scikit learn `linearregression`.

✓ Part 1: Method of least squares

```
X0=np.ones(422)
X1=np.transpose(diabetes_X_train)
X=np.vstack((X0,X1))
Xtrans=np.transpose(X)
product1=np.matmul(X , diabetes_y_train)
product2= np.matmul(X , Xtrans)
product2inv= np.linalg.inv(product2)
theta= np.matmul(product2inv , product1)
```

```
print ("theta0 :",theta[0])
print ("theta1 :",theta[1])
```

```
theta0 : 152.91886182616113
theta1 : 938.2378612513521
```

▼ Part 2 : Linear Regression

```
epochs=11000
lr=0.45
theta=np.zeros(2) #initialize weights to be 0
for i in range (epochs):
    output=np.dot(Xtrans,theta)
    error=output-diabetes_y_train
    loss=np.sum(error**2)/(2*422) #422 is number of test data rows
    gradient=np.dot(X,error)/422
    theta=theta-lr*gradient

print (theta)
```

→ [152.91886781 938.22523304]

Part 3: Gradient for new Loss Function

$$L = \sum_{i=1}^n \frac{1}{2} (y^{(i)} - h_{\theta}(x^{(i)}))^2 + \frac{\lambda}{2} \| \theta \|^2_2, \text{ where } \theta = [\theta_0 \ \theta_1]^T$$

$$= \sum_{i=1}^n \frac{1}{2} (y^{(i)} - h_{\theta}(x^{(i)}))^2 + \frac{\lambda}{2} (\theta_1^2 + \theta_0^2)$$

$$\nabla L_{\theta} = \begin{bmatrix} \frac{d}{d\theta_0} \left(\frac{1}{2} (y^{(i)} - h_{\theta}(x^{(i)}))^2 + \frac{\lambda}{2} (\theta_1^2 + \theta_0^2) \right) \\ \frac{d}{d\theta_1} \left(\frac{1}{2} (y^{(i)} - h_{\theta}(x^{(i)}))^2 + \frac{\lambda}{2} (\theta_1^2 + \theta_0^2) \right) \end{bmatrix}$$

$$= \begin{bmatrix} \frac{d}{d\theta_0} \left(\frac{1}{2} (y^{(i)} - h_{\theta}(\theta_0 + \theta_1 x^{(i)}))^2 + \frac{\lambda}{2} (\theta_1^2 + \theta_0^2) \right) \\ \frac{d}{d\theta_1} \left(\frac{1}{2} (y^{(i)} - h_{\theta}(\theta_0 + \theta_1 x^{(i)}))^2 + \frac{\lambda}{2} (\theta_1^2 + \theta_0^2) \right) \end{bmatrix}$$

$$= \begin{bmatrix} -(y^{(i)} - (\theta_0 + \theta_1 x^{(i)})) + \lambda \theta_0 \\ -(y^{(i)} - (\theta_0 + \theta_1 x^{(i)})) x^{(i)} + \lambda \theta_1 \end{bmatrix}$$

$$= \begin{bmatrix} \lambda \theta_0 + \theta_0 + \theta_1 x^{(i)} - y^{(i)} \\ \lambda \theta_1 + (x^{(i)}) (\theta_0 + \theta_1 x^{(i)}) - y^{(i)} \end{bmatrix}$$